

RedisTimeSeries

SILICON VALLEY ENGINEER

- RedisTimeSeries is a module **designed to efficiently store and query time series data within Redis**.
- It provides commands for creating time series, updating them with new observations, and querying data over time ranges.
- RedisTimeSeries is optimized for **time series data**, supporting downsampling, aggregation, and retention policies.
- This module is ideal for IoT, analytics, and monitoring applications where time-stamped data is continuously generated and queried.



RedisTimeSeries - Feature

SILICON VALLEY ENGINEER

- **Time-based Queries:** Enables querying of data based on time ranges, simplifying the retrieval of historical data within specific intervals.
- **Aggregations:** Provides built-in aggregation functions (e.g., sum, avg, min, max, count) over time windows, allowing for the easy computation of metrics over time intervals directly on the database.
- **Retention Policies:** Enables the configuration of data retention policies, automatically expiring and removing old data to manage disk space usage efficiently.
- **Labels and Filtering:** Allows the assignment of labels to time series and supports filtering queries by these labels, facilitating the organization and querying of data from different sources or categories.
- **Downsampling:** Supports automatic downsampling of data, which is critical for long-term data retention strategies and reducing the resolution of data as it ages.